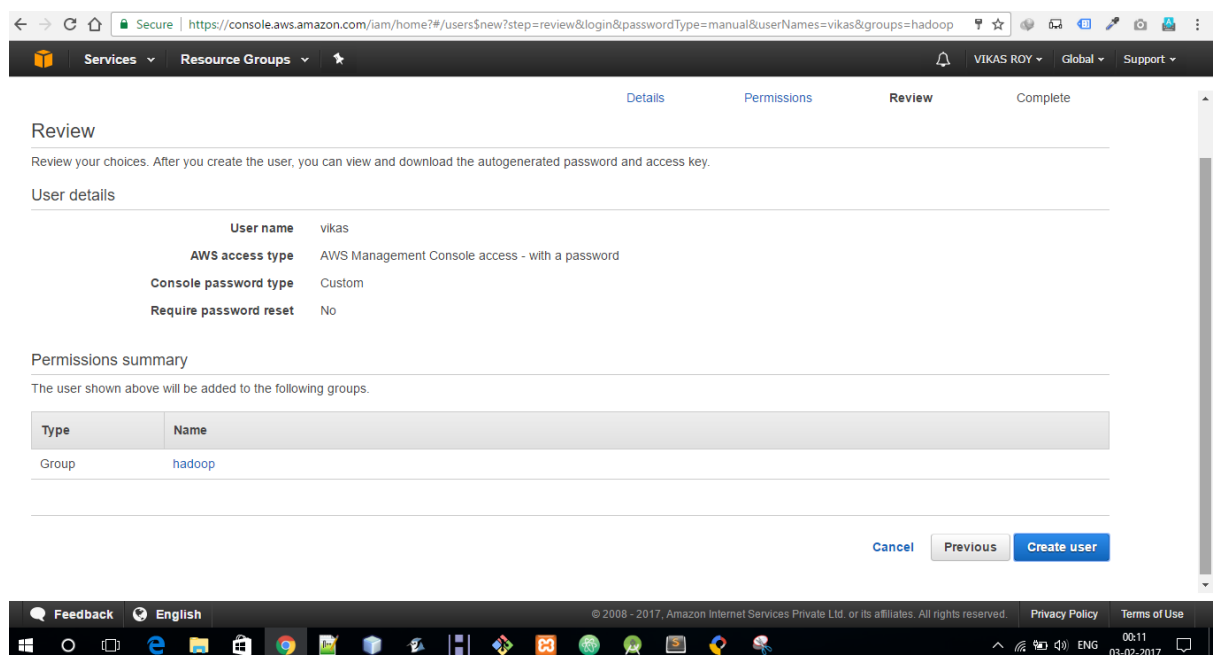
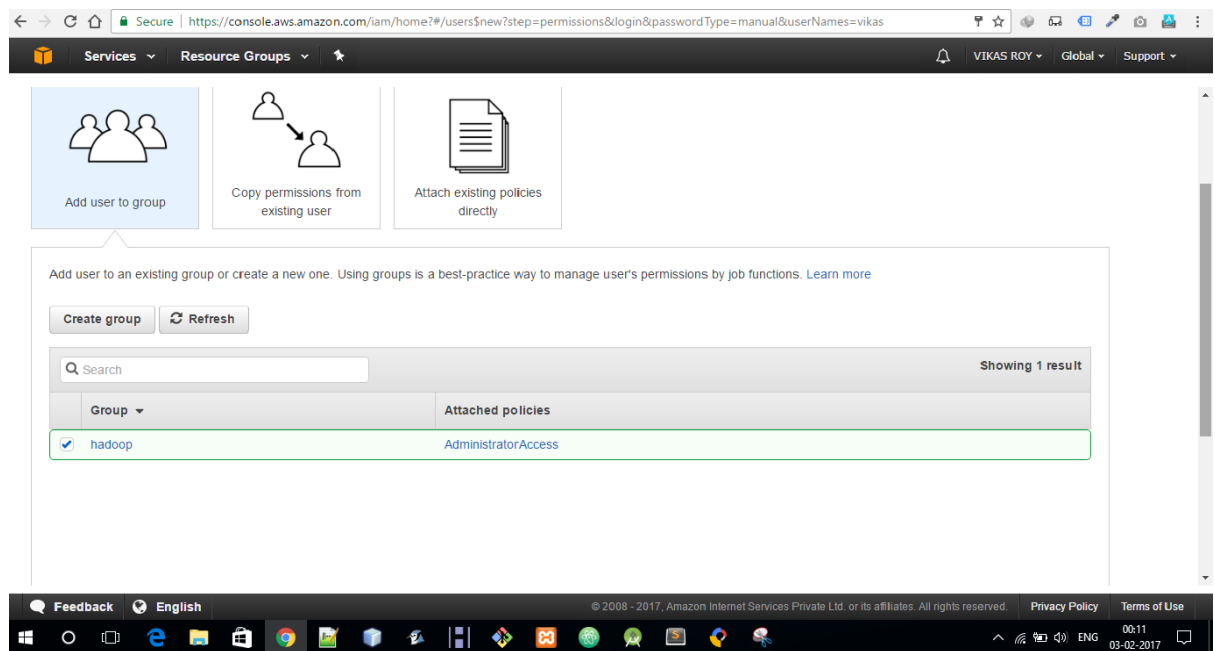
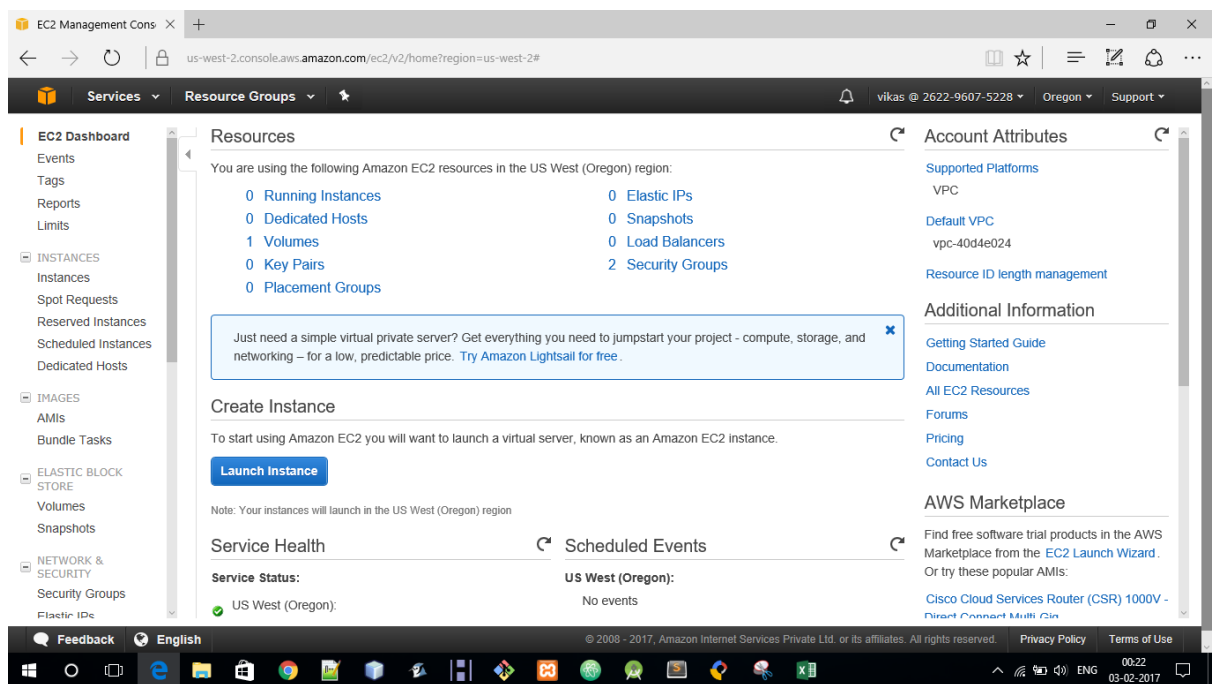
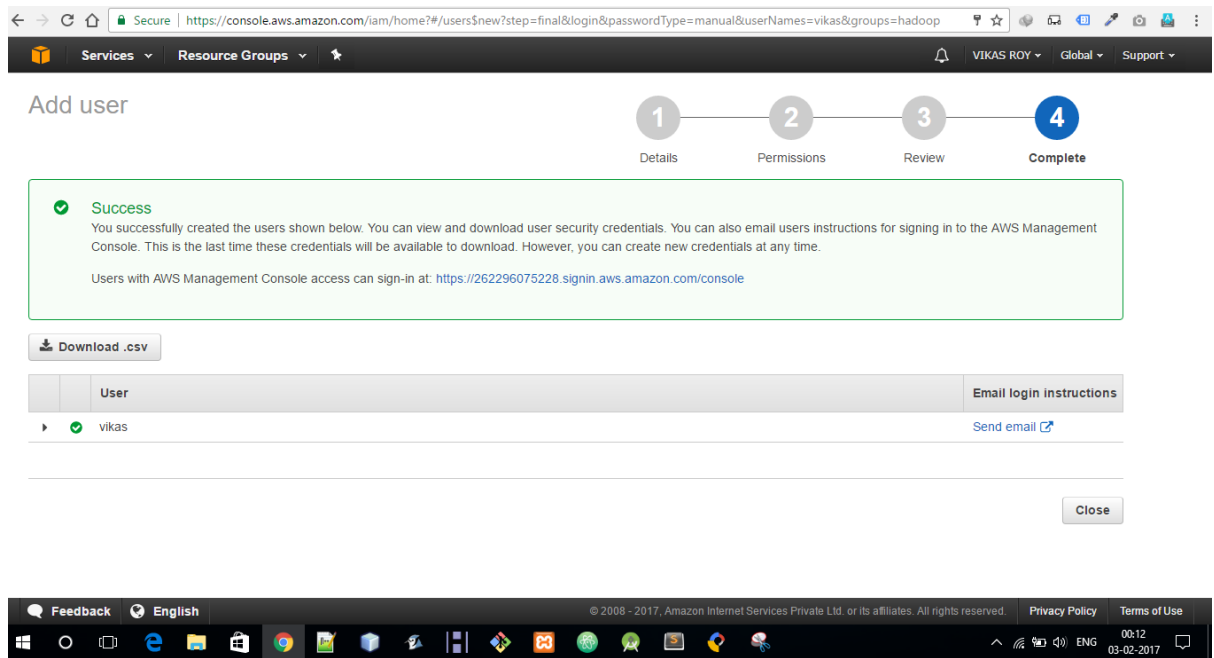


How to install Hadoop in Aws EC2

Before installing Hadoop we have to Create instance in EC2

1. First create account at aws.amazon.com
2. After that create an IAM user via <https://console.aws.amazon.com/iam/>
3. Add user to group
4. Create Group and give admin access to that group

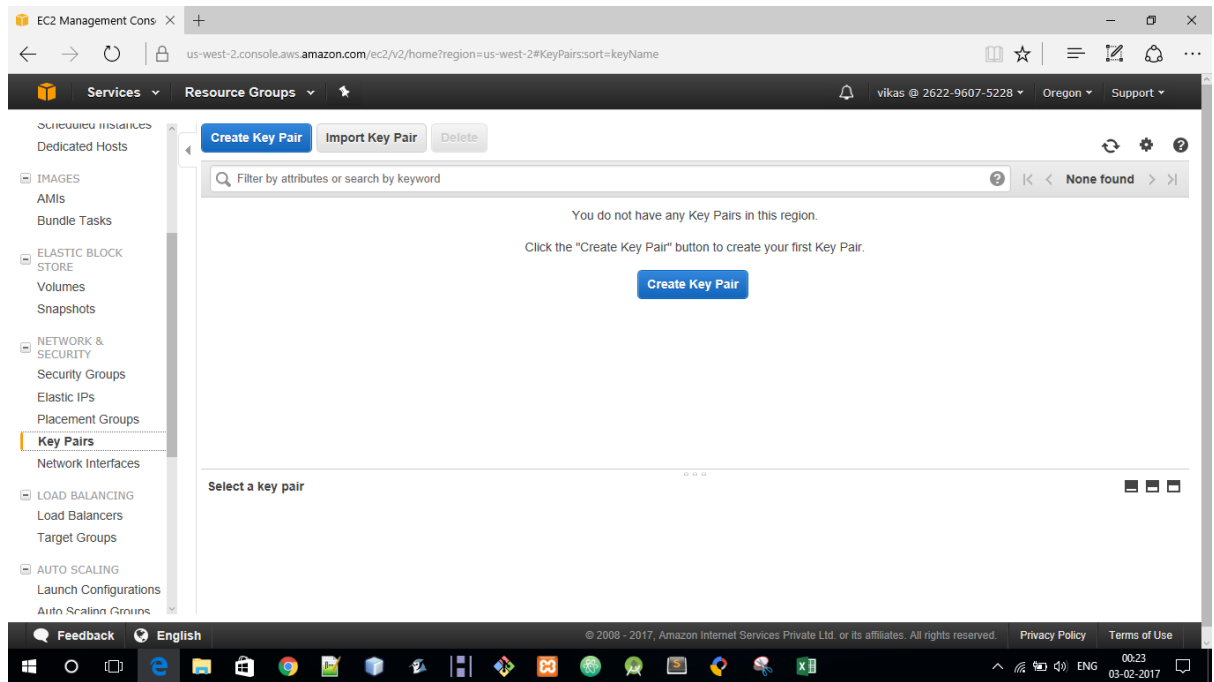




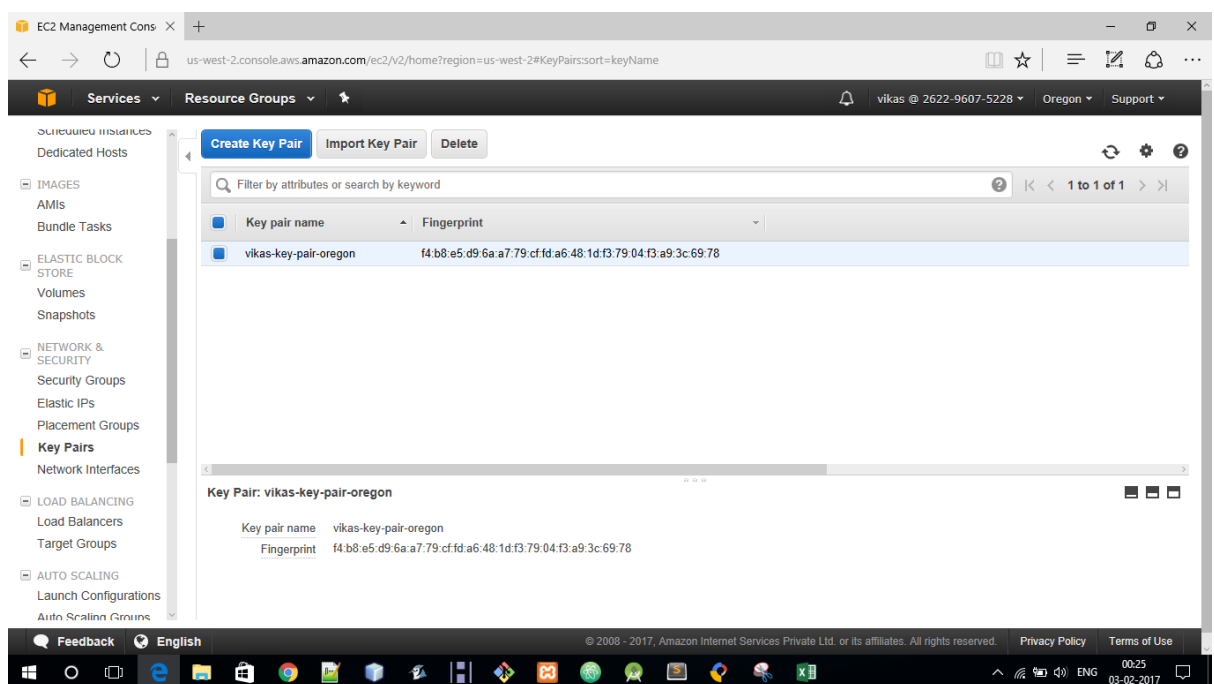
To sign in as IAM new user we need aws account id in following url and then it can be login by providing user name (this is not same as aws account email) and password.

Create a Key Pair

1. Sign in to aws using above url
2. From dashboard choose EC2 to open the amazon EC2 console
3. In navigation pane under Network and security , click key pairs

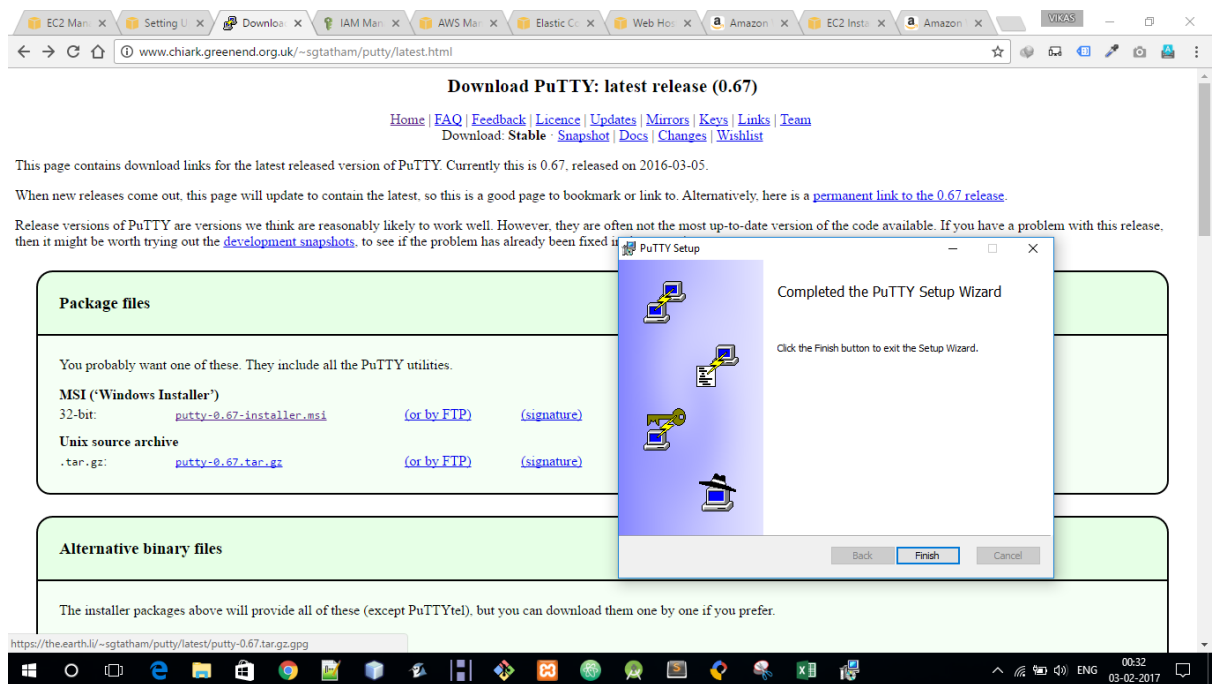
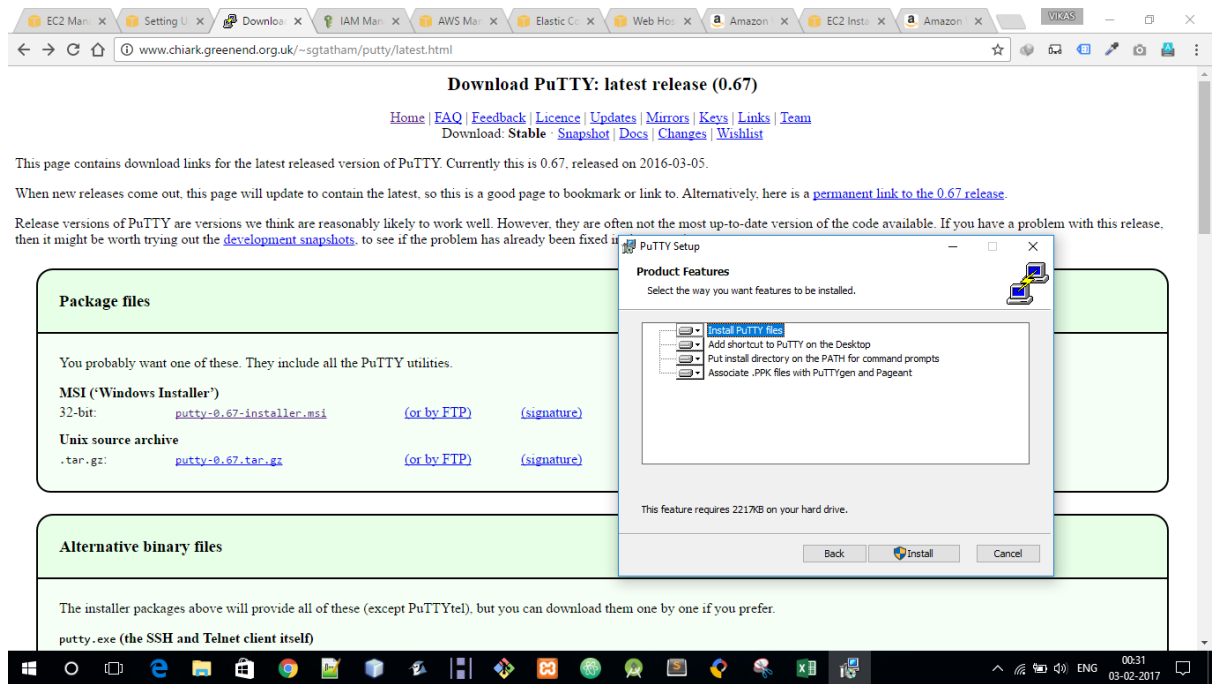


4. Click Create key pair
5. Enter name for new key pair followed by `-key-pair`

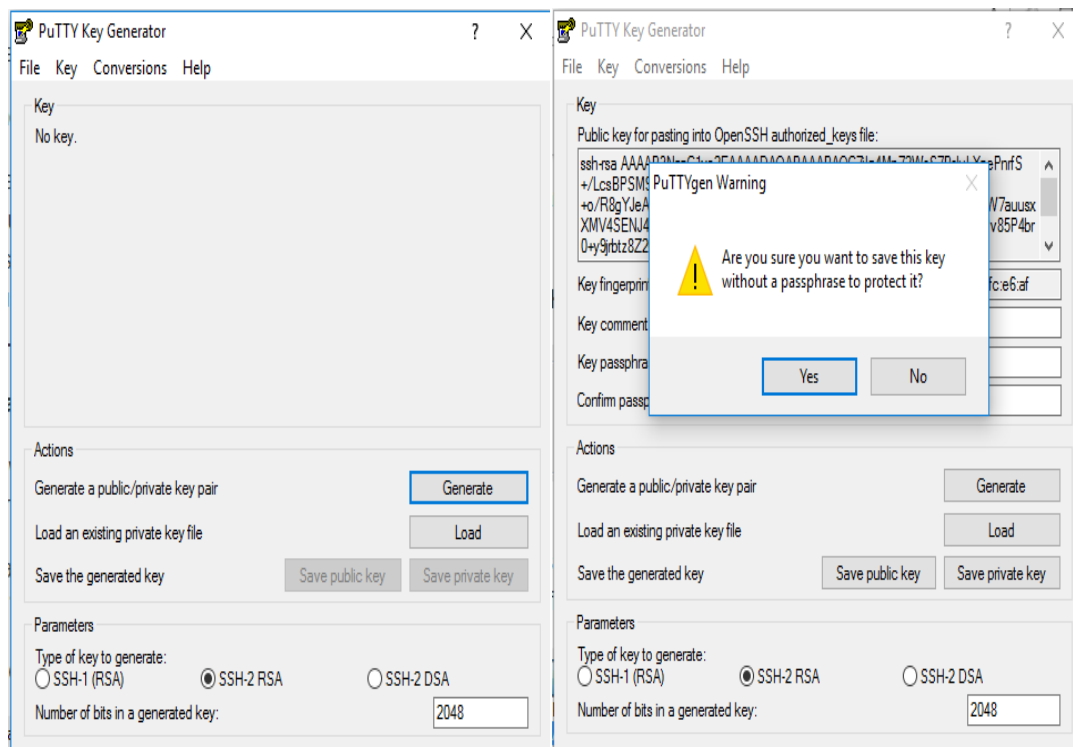


6. Private key will be downloaded automatically with .pem extension

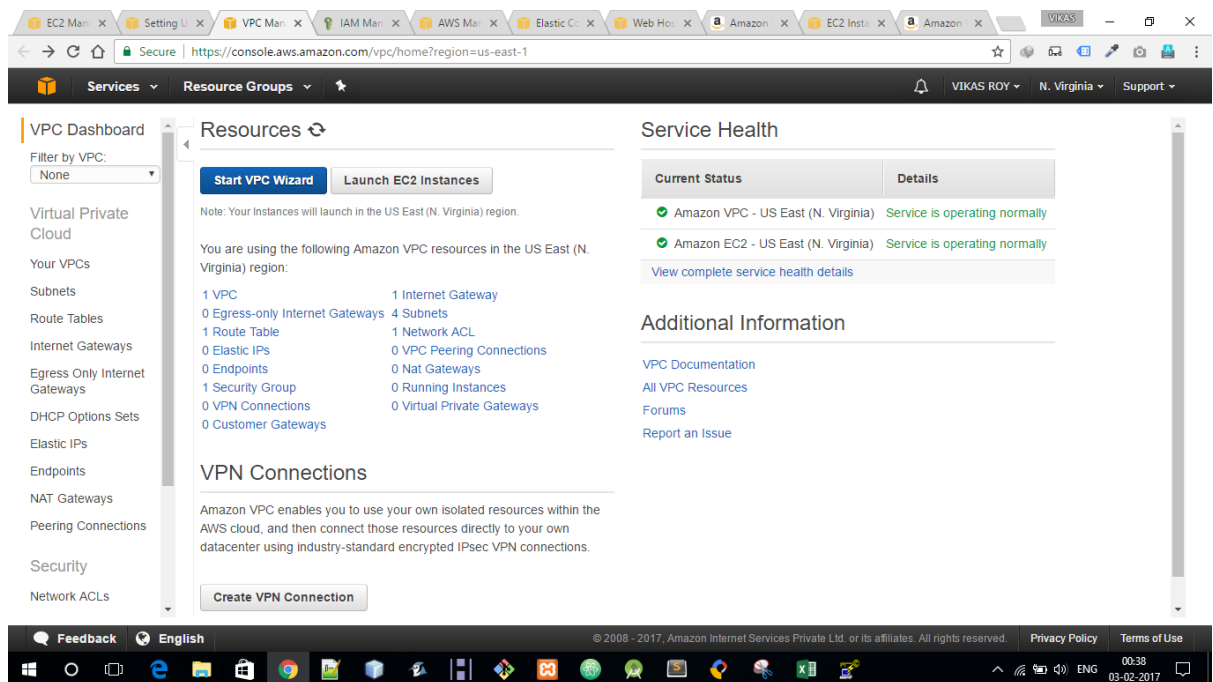
For connection with linux instance via windows or linux we can download putty entire package.



Using putty keygen we have to convert our .pem file which we have downloaded earlier into .ppk file



If by default VPC is not created we have to create VPC via VPC dashboard



EC2 Man X Setting U X VPC Man X IAM Man X AWS Man X Elastic C X Web Ho X Amazon X EC2 Insta X Amazon X VIKAS

Secure | https://console.aws.amazon.com/vpc/home?region=us-east-1#wizardSelector:

Services Resource Groups

VIKAS ROY N. Virginia Support

Step 1: Select a VPC Configuration

VPC with a Single Public Subnet

VPC with Public and Private Subnets

VPC with Public and Private Subnets and Hardware VPN Access

VPC with a Private Subnet Only and Hardware VPN Access

Your instances run in a private, isolated section of the AWS cloud with direct access to the Internet. Network access control lists and security groups can be used to provide strict control over inbound and outbound network traffic to your instances.

Creates:

A /16 network with a /24 subnet. Public subnet instances use Elastic IPs or Public IPs to access the Internet.

Select

Internet, S3, DynamoDB, SNS, SQS, etc.

Public Subnet

Amazon Virtual Private Cloud

Cancel and Exit

Feedback English © 2008 - 2017, Amazon Internet Services Private Ltd. or its affiliates. All rights reserved. Privacy Policy Terms of Use

EC2 Man X Setting U X VPC Man X IAM Man X AWS Man X Elastic C X Web Ho X Amazon X EC2 Insta X Amazon X VIKAS

Secure | https://console.aws.amazon.com/vpc/home?region=us-east-1#wizardFullpagePublicOnly:

Services Resource Groups

VIKAS ROY N. Virginia Support

Step 2: VPC with a Single Public Subnet

IPv4 CIDR block:* 10.0.0.0/16 (65531 IP addresses available)

IPv6 CIDR block: ☒ No IPv6 CIDR Block ☐ Amazon provided IPv6 CIDR block

VPC name: vpc

Public subnet's IPv4 CIDR:* 10.0.0.0/24 (251 IP addresses available)

Availability Zone:* No Preference

Subnet name: Public subnet

You can add more subnets after AWS creates

Service endpoints

Add Endpoint

Enable DNS hostnames:* ☒ Yes ☐ No

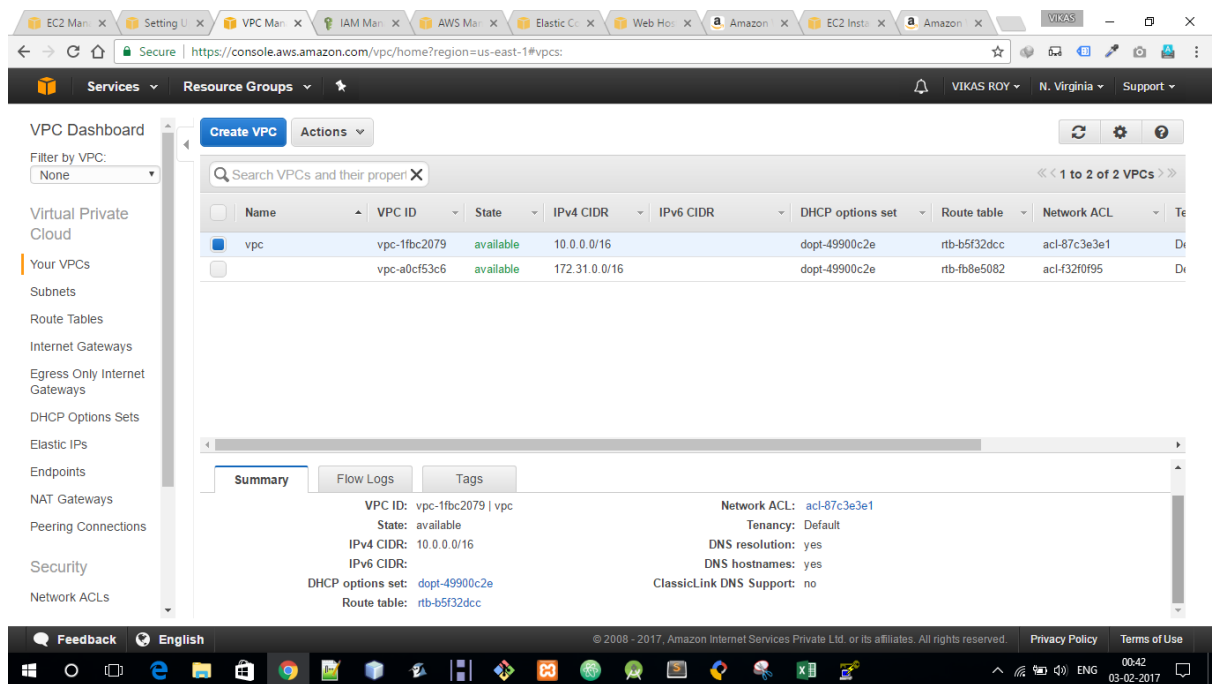
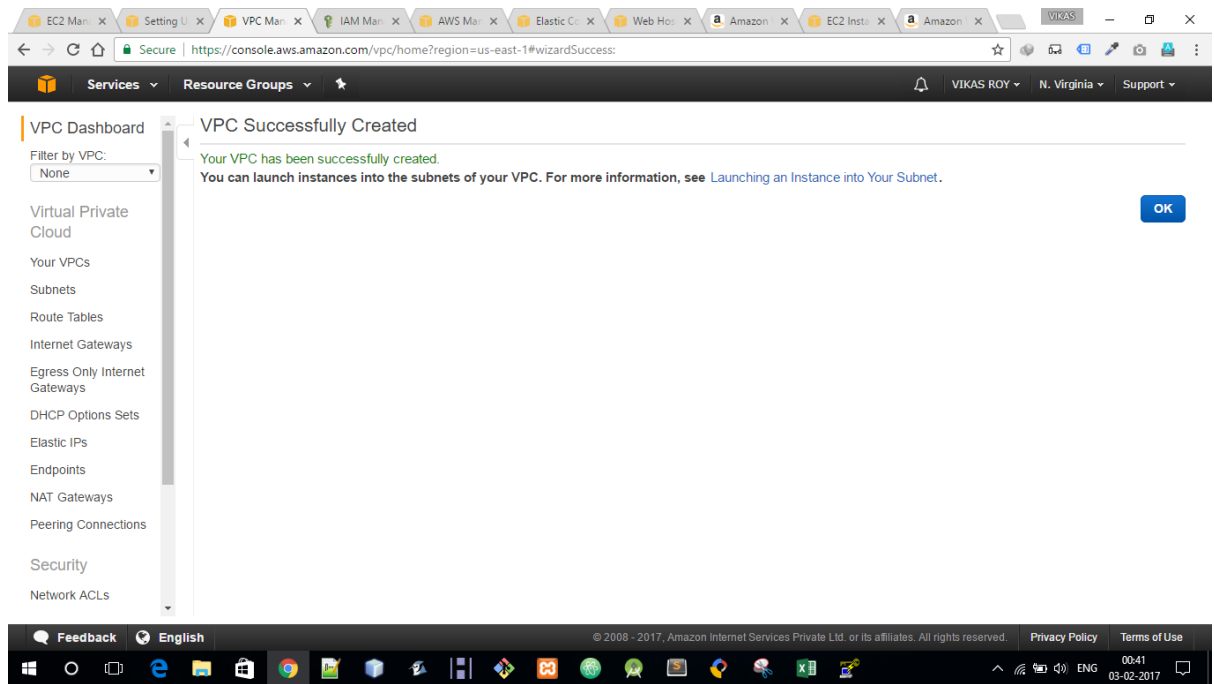
Hardware tenancy:* Default

Cancel and Exit Back Create VPC

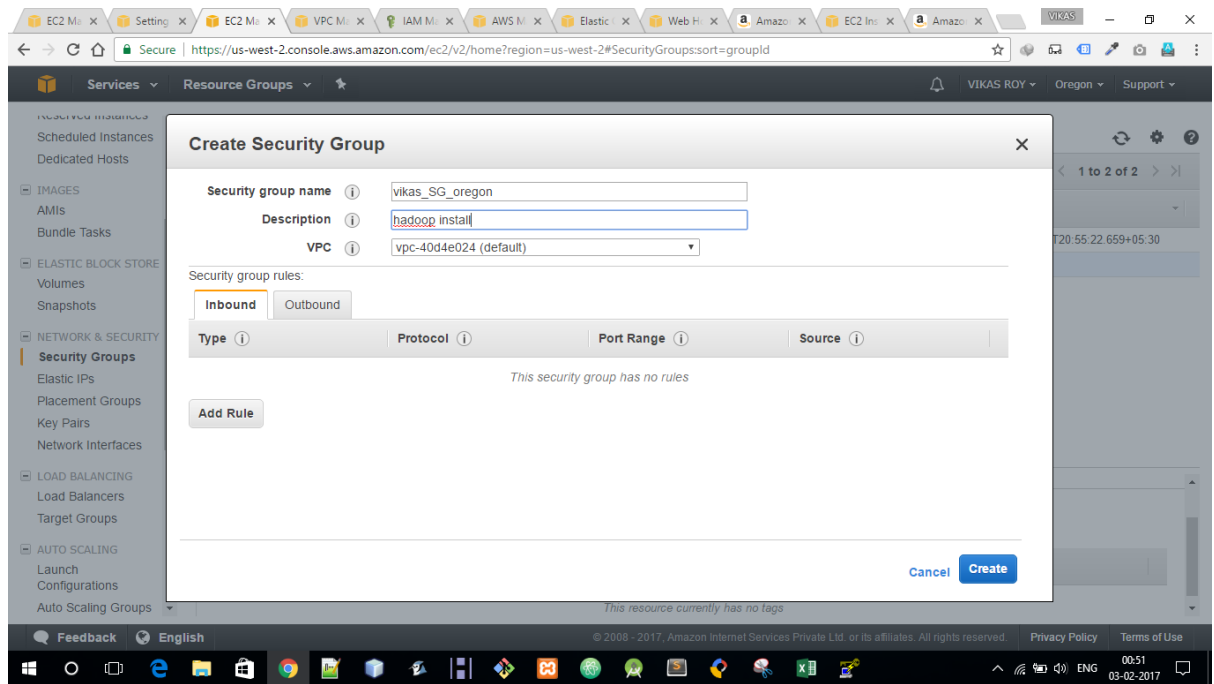
Associating Subnet to Route Table...

33%

Feedback English © 2008 - 2017, Amazon Internet Services Private Ltd. or its affiliates. All rights reserved. Privacy Policy Terms of Use



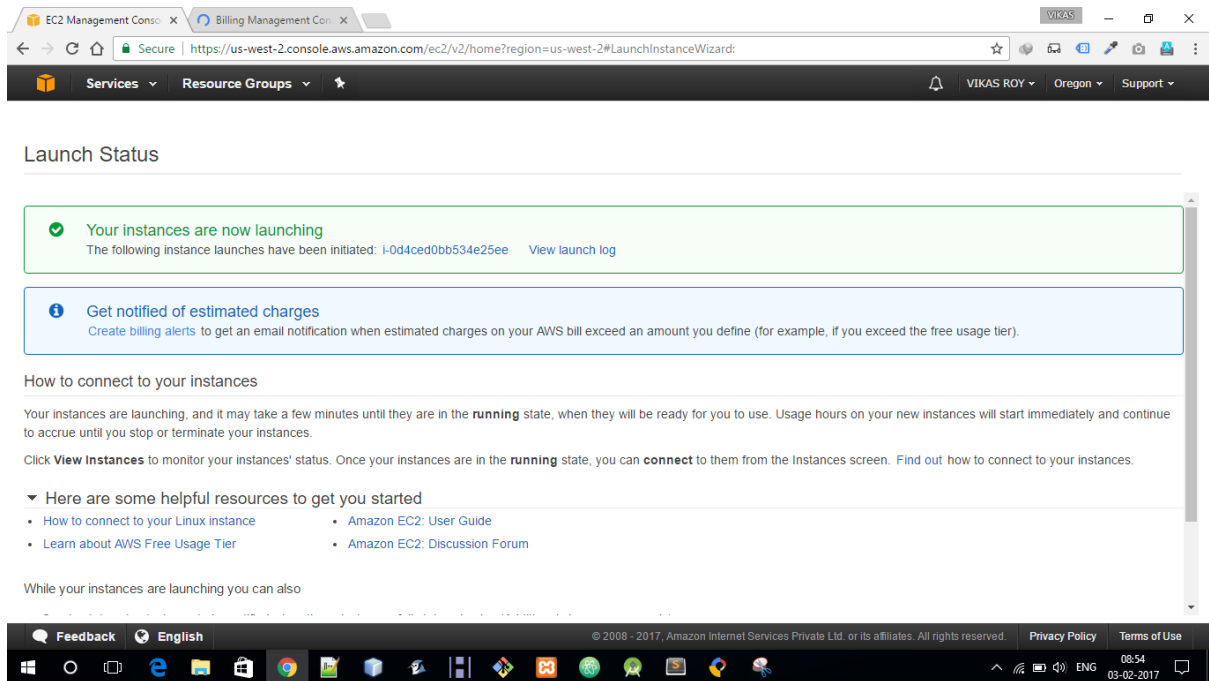
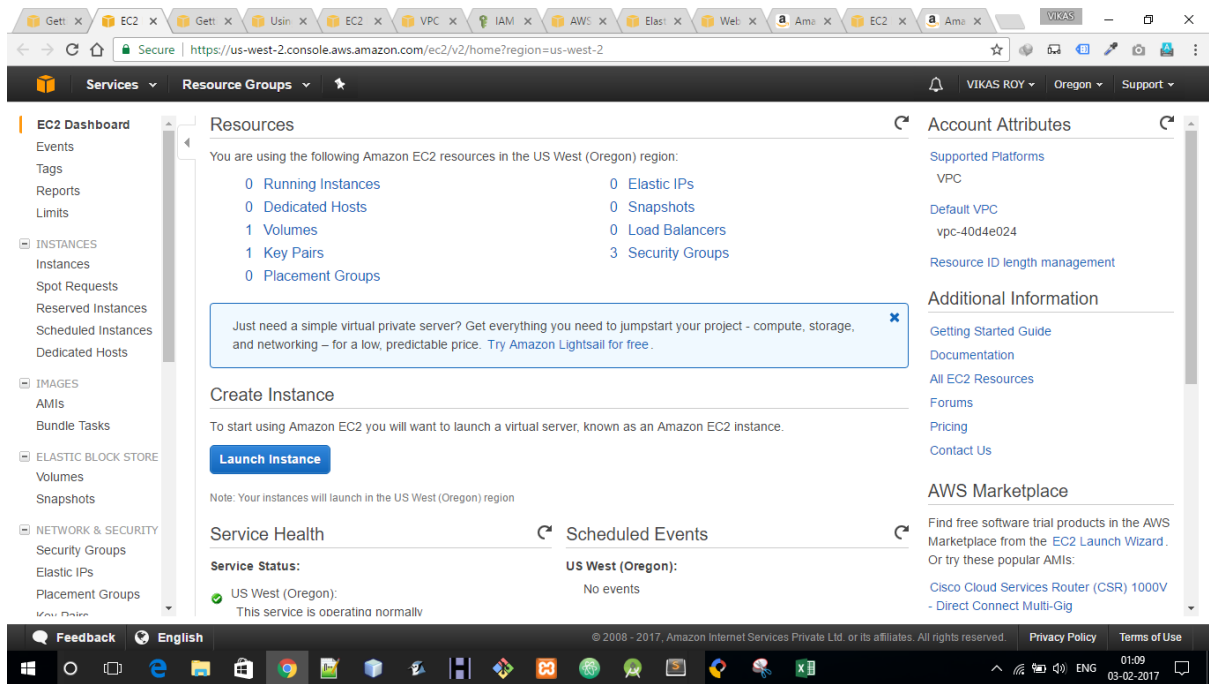
After that create security group by accessing navigation pane.



1. Click security group in navigation pane
2. Click create security group
3. Name security group, description and select vpc default list from drop down.
4. In Inbound tab create the following rule
 - Select HTTP from the Type list, and make sure that Source is set to Anywhere
 - Select HTTPS from the Type list, and make sure that Source is set to Anywhere
 - Select SSH from the Type list. In the Source box, choose My IP to automatically populate the field with the public IPv4 address of your local computer.

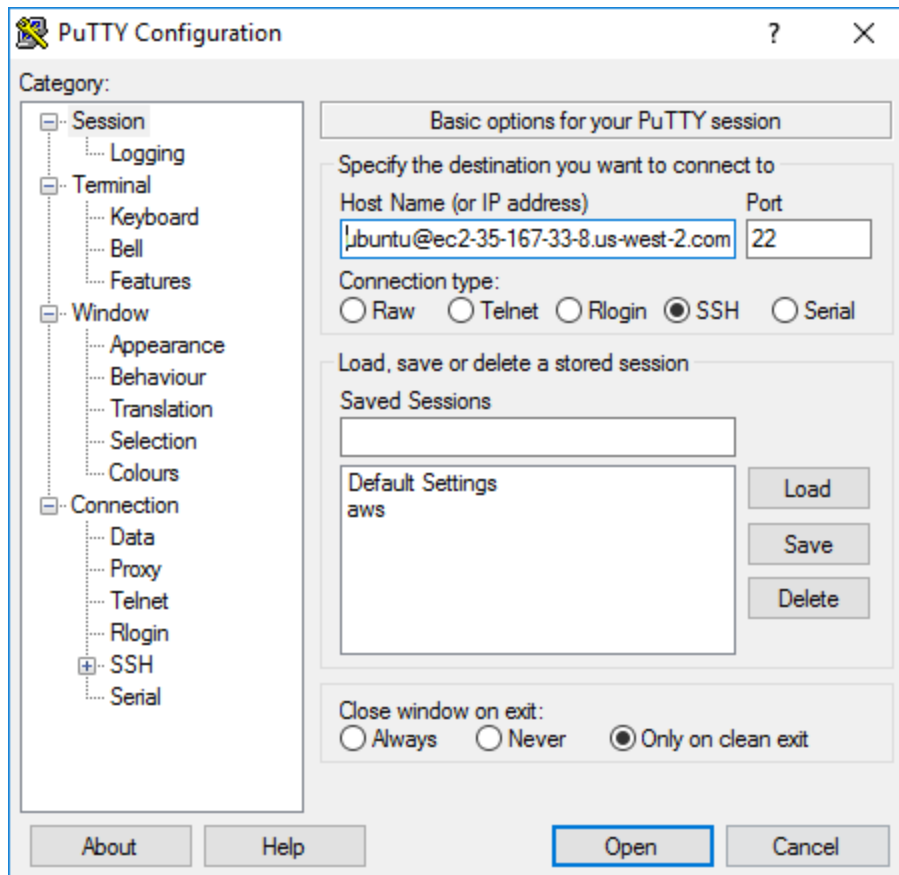
Launch instance

Select os which you want to launch , and select existing security group when it ask for it , choose existing key pair. Then it will take some time and instance will be launch automitacally.




Connect instance from windows using putty

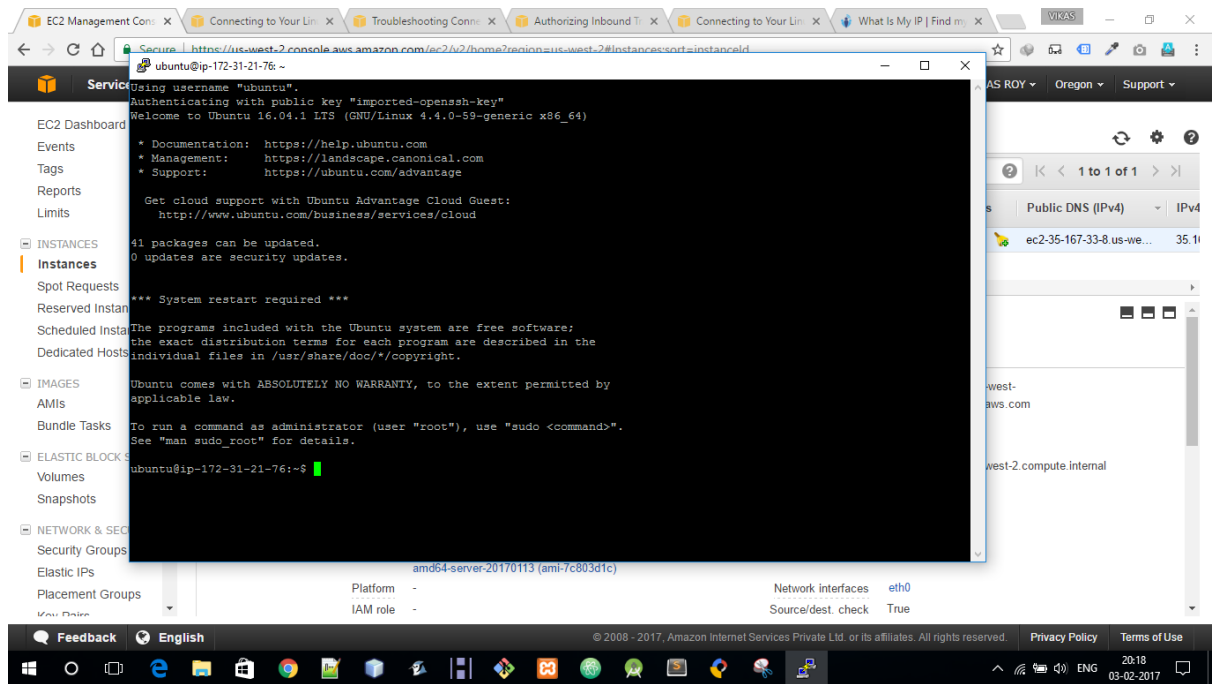
1. Start putty
2. In host name box enter username@public-dns
3. In my case I have Ubuntu running instance so my username will be Ubuntu and public dns can be seen from aws dashboard



4. Under Connection expand ssh and select auth choose browse and select ppk file which you have generated earlier. Then Click Open.
5. During first time connection putty will display security alert ignore and select yes.



6. A window will be open and that will be connected to instance.



How to install Hadoop

Before starting Hadoop installation we have to check whether java is install or not in our system

Type following command in instance for check

Java -version

```
ubuntu@ip-172-31-21-76: ~
* Documentation: https://help.ubuntu.com
* Management: https://landscape.canonical.com
* Support: https://ubuntu.com/advantage

Get cloud support with Ubuntu Advantage Cloud Guest:
http://www.ubuntu.com/business/services/cloud

41 packages can be updated.
0 updates are security updates.

*** System restart required ***

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

ubuntu@ip-172-31-21-76:~$ java -version
The program 'java' can be found in the following packages:
* default-jre
* gcj-5-jre-headless
* openjdk-8-jre-headless
* gcj-4.8-jre-headless
* gcj-4.9-jre-headless
* openjdk-9-jre-headless
Try: sudo apt install <selected package>
ubuntu@ip-172-31-21-76:~$
```

If java is not installed try “sudo apt install openjdk-9-jre-headless”

You can install different version of java jdk also, I have downloaded latest one in my case.

```
ubuntu@ip-172-31-21-76: ~  
* gcj-5-jre-headless  
* openjdk-8-jre-headless  
* gcj-4.8-jre-headless  
* gcj-4.9-jre-headless  
* openjdk-9-jre-headless  
Try: sudo apt install <selected package>  
ubuntu@ip-172-31-21-76:~$ ls  
ubuntu@ip-172-31-21-76:~$ java --version  
The program 'java' can be found in the following packages:  
* default-jre  
* gcj-5-jre-headless  
* openjdk-8-jre-headless  
* gcj-4.8-jre-headless  
* gcj-4.9-jre-headless  
* openjdk-9-jre-headless  
Try: sudo apt install <selected package>  
ubuntu@ip-172-31-21-76:~$ sudo apt install openjdk-9-jre-headless  
Reading package lists... Done  
Building dependency tree  
Reading state information... Done  
The following additional packages will be installed:  
  ca-certificates-java java-common libavahi-client3 libavahi-common-data libavahi-common3 libcups2 libjpeg-turbo8  
  libjpeg8 liblcms2-2 libnsspr4 libnss3 libnss3-nssdb libpcsclite1 libxi6 libxrender1 libxtst6 x11-common  
Suggested packages:  
  default-jre cups-common liblcms2-utils pcsd libnss-mdns fonts-dejavu-extra fonts-ipafont-gothic  
  fonts-ipafont-mincho ttf-wqy-microhei | ttf-wqy-zenhei fonts-indic  
The following NEW packages will be installed:  
  ca-certificates-java java-common libavahi-client3 libavahi-common-data libavahi-common3 libcups2 libjpeg-turbo8  
  libjpeg8 liblcms2-2 libnsspr4 libnss3 libnss3-nssdb libpcsclite1 libxi6 libxrender1 libxtst6 openjdk-9-jre-headless  
  x11-common  
0 upgraded, 18 newly installed, 0 to remove and 46 not upgraded.  
Need to get 185 MB of archives.  
After this operation, 313 MB of additional disk space will be used.  
Do you want to continue? [Y/n]
```


```
ubuntu@ip-172-31-21-76: ~  
Adding debian:Atos_TrustedRoot_2011.pem  
Adding debian:TWCA_Global_Root_CA.pem  
Adding debian:TeliaSonera_Root_CA_v1.pem  
Adding debian:DigiCert_Assured_ID_Root_CA.pem  
Adding debian:WoSign.pem  
Adding debian:AddTrust_Qualified_Certificates_Root.pem  
Adding debian:GeoTrust_Primary_Certification_Authority.pem  
Adding debian:CA_Disig.pem  
Adding debian:QuoVadis_Root_CA_2_G3.pem  
Adding debian:SwissSign_Gold_CA_-_G2.pem  
Adding debian:Sonera_Class_2_Root_CA.pem  
Adding debian:Security_Communication_RootCA2.pem  
Adding debian:Verisign_Class_2_Public_Primary_Certification_Authority_-_G2.pem  
Adding debian:Certigna.pem  
Adding debian:D-TRUST_Root_Class_3_CA_2_EV_2009.pem  
Adding debian:Verisign_Class_3_Public_Primary_Certification_Authority_-_G2.pem  
Adding debian:COMODO_ECC_Certification_Authority.pem  
Adding debian:OISTE_WISKey_Global_Root_GB_CA.pem  
Adding debian:CFCA_EV_ROOT.pem  
Adding debian:Comodo_Trusted_Services_root.pem  
done.  
Processing triggers for libc-bin (2.23-0ubuntu5) ...  
Processing triggers for systemd (229-4ubuntu13) ...  
Processing triggers for ureadahead (0.100.0-19) ...  
Processing triggers for ca-certificates (20160104ubuntu1) ...  
Updating certificates in /etc/ssl/certs...  
0 added, 0 removed; done.  
Running hooks in /etc/ca-certificates/update.d...  
  
done.  
done.  
ubuntu@ip-172-31-21-76:~$
```

After successful installation of java find mirror of Hadoop and then download it in aws linux instance using **wget** command

Apache Download Mirror: X

www.apache.org/dyn/closer.cgi/hadoop/common/hadoop-2.6.0/hadoop-2.6.0.tar.gz

Home » Dyn About Projects People Get Involved Download Support Apache

 THE APACHE SOFTWARE FOUNDATION

We suggest the following mirror site for your download:
<http://www-us.apache.org/dist/hadoop/common/hadoop-2.6.0/hadoop-2.6.0.tar.gz>

Other mirror sites are suggested below. Please use the backup mirrors only to download PGP and MD5 signatures to verify your downloads or if no other mirrors are working.

HTTP

<http://www-eu.apache.org/dist/hadoop/common/hadoop-2.6.0/hadoop-2.6.0.tar.gz>

<http://www-us.apache.org/dist/hadoop/common/hadoop-2.6.0/hadoop-2.6.0.tar.gz>

BACKUP SITES

Please use the backup mirrors only to download PGP and MD5 signatures to verify your downloads or if no other mirrors are working.

<http://www-eu.apache.org/dist/hadoop/common/hadoop-2.6.0/hadoop-2.6.0.tar.gz>

<http://www-us.apache.org/dist/hadoop/common/hadoop-2.6.0/hadoop-2.6.0.tar.gz>

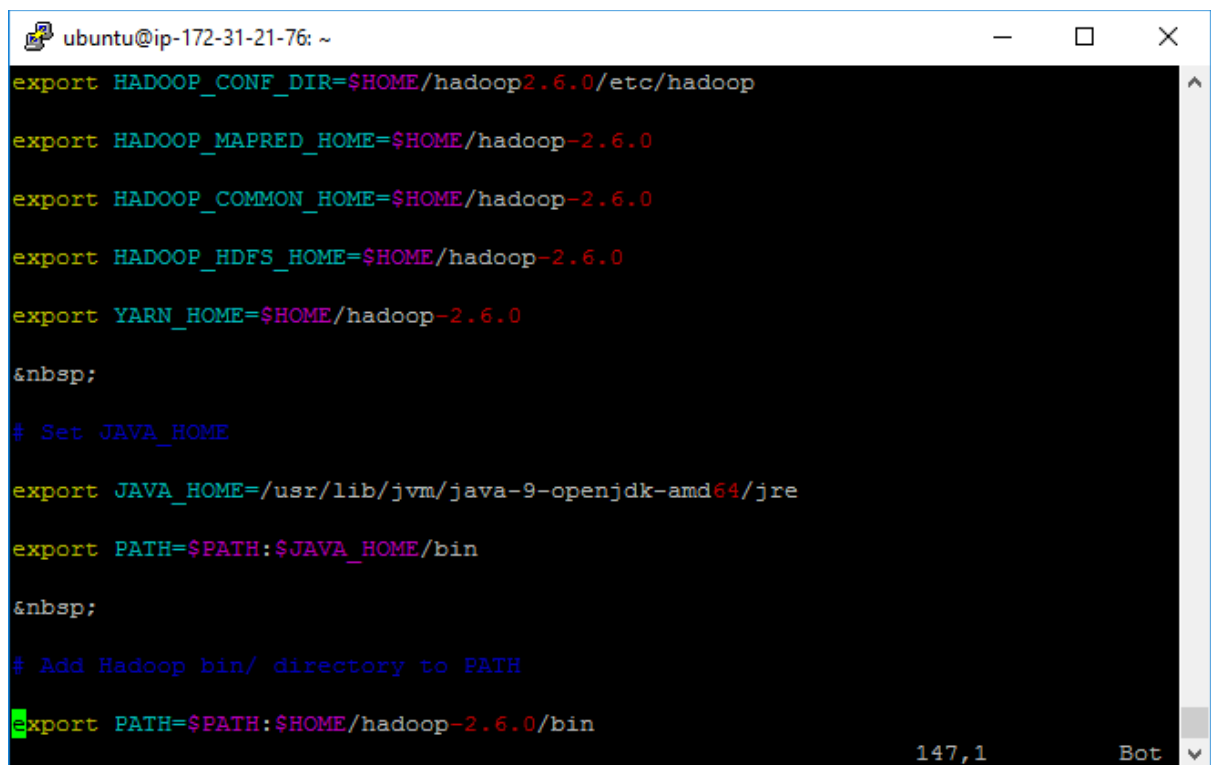
```
ubuntu@ip-172-31-21-76: ~  
Adding debian:Sonera_Class_2_Root_CA.pem  
Adding debian:Security_Communication_RootCA2.pem  
Adding debian:Verisign_Class_2_Public_Primary_Certification_Authority_-_G2.pem  
Adding debian:Certigna.pem  
Adding debian:D-TRUST_Root_Class_3_CA_2_EV_2009.pem  
Adding debian:Verisign_Class_3_Public_Primary_Certification_Authority_-_G2.pem  
Adding debian:COMODO_ECC_Certification_Authority.pem  
Adding debian:OISeKey_Global_Root_GB_CA.pem  
Adding debian:CFCA_EV_ROOT.pem  
Adding debian:Comodo_Trusted_Services_root.pem  
done.  
Processing triggers for libc-bin (2.23-0ubuntu5) ...  
Processing triggers for systemd (229-4ubuntu13) ...  
Processing triggers for ureadahead (0.100.0-19) ...  
Processing triggers for ca-certificates (20160104ubuntu1) ...  
Updating certificates in /etc/ssl/certs...  
0 added, 0 removed; done.  
Running hooks in /etc/ca-certificates/update.d...  
  
done.  
done.  
ubuntu@ip-172-31-21-76:~$ java -version  
openjdk version "9-internal"  
OpenJDK Runtime Environment (build 9-internal+0-2016-04-14-195246.builddd.src)  
OpenJDK 64-Bit Server VM (build 9-internal+0-2016-04-14-195246.builddd.src, mixed mode)  
ubuntu@ip-172-31-21-76:~$ wget http://www.trieuvan.com/apache/hadoop/common/hadoop-2.6.0/hadoop-2.6.0.tar.gz  
--2017-02-03 15:03:15-- http://www.trieuvan.com/apache/hadoop/common/hadoop-2.6.0/hadoop-2.6.0.tar.gz  
Resolving www.trieuvan.com (www.trieuvan.com)... 66.201.46.168  
Connecting to www.trieuvan.com (www.trieuvan.com)|66.201.46.168|:80... connected.  
HTTP request sent, awaiting response... 200 OK  
Length: 195257604 (186M) [application/x-gzip]  
Saving to: 'hadoop-2.6.0.tar.gz'  
  
hadoop-2.6.0.tar.gz      28%[=====>] 52.62M  10.8MB/s  eta 13s
```



```
"# Set Hadoop-related environment variables
export HADOOP_HOME=$HOME/hadoop-2.6.0
export HADOOP_CONF_DIR=$HOME/hadoop2.6.0/etc/hadoop
export HADOOP_MAPRED_HOME=$HOME/hadoop-2.6.0
export HADOOP_COMMON_HOME=$HOME/hadoop-2.6.0
export HADOOP_HDFS_HOME=$HOME/hadoop-2.6.0
export YARN_HOME=$HOME/hadoop-2.6.0

# Set JAVA_HOME
export JAVA_HOME=/usr/lib/jvm/java-9-openjdk-amd64/bin/java
export PATH=$PATH:$JAVA_HOME/bin

# Add Hadoop bin/ directory to PATH
export PATH=$PATH:$HOME/hadoop-2.6.0/bin"
```

A terminal window titled 'ubuntu@ip-172-31-21-76: ~' with standard window controls. The terminal displays the same environment variable configuration code as the first block. The text is color-coded: '#' is blue, 'export' is green, and paths are in pink. The code is as follows:

```
export HADOOP_CONF_DIR=$HOME/hadoop2.6.0/etc/hadoop
export HADOOP_MAPRED_HOME=$HOME/hadoop-2.6.0
export HADOOP_COMMON_HOME=$HOME/hadoop-2.6.0
export HADOOP_HDFS_HOME=$HOME/hadoop-2.6.0
export YARN_HOME=$HOME/hadoop-2.6.0
&nbsp;
# Set JAVA_HOME
export JAVA_HOME=/usr/lib/jvm/java-9-openjdk-amd64/jre
export PATH=$PATH:$JAVA_HOME/bin
&nbsp;
# Add Hadoop bin/ directory to PATH
export PATH=$PATH:$HOME/hadoop-2.6.0/bin
```

The bottom right corner of the terminal shows '147,1' and 'Bot' with a dropdown arrow.

Use "source .bashrc" for refresh and check .bashrc file

Then enter Hadoop for check environment is working or not

```
ubuntu@ec2-user: ~
-bash: .bashrc: line 142: `&nbsp;'
ubuntu@ec2-user:~$ sudo vi .bashrc
ubuntu@ec2-user:~$ source .bashrc
ubuntu@ec2-user:~$ hadoop
Usage: hadoop [--config confdir] COMMAND
      where COMMAND is one of:
      fs                run a generic filesystem user client
      version           print the version
      jar <jar>         run a jar file
      checknative [-a|-h] check native hadoop and compression libraries availability
      distcp <srcurl> <desturl> copy file or directories recursively
      archive -archiveName NAME -p <parent path> <src>* <dest> create a hadoop archive
      classpath         prints the class path needed to get the
      credential        interact with credential providers
                        Hadoop jar and the required libraries
      daemonlog         get/set the log level for each daemon
      trace             view and modify Hadoop tracing settings
      or
      CLASSNAME         run the class named CLASSNAME

Most commands print help when invoked w/o parameters.
ubuntu@ec2-user:~$
```

- Cd Hadoop-2.6.0
- Cd share/Hadoop

```
ubuntu@ec2-user: ~/hadoop-2.6.0/sbin
ubuntu@ec2-user:~$ cd hadoop-2.6.0
ubuntu@ec2-user:~/hadoop-2.6.0$ cd share/hadoop
ubuntu@ec2-user:~/hadoop-2.6.0/share/hadoop$ ls
common  hdfs  httpfs  kms  mapreduce  tools  yarn
ubuntu@ec2-user:~/hadoop-2.6.0/share/hadoop$ cd ..
ubuntu@ec2-user:~/hadoop-2.6.0/share$ ls
doc  hadoop
ubuntu@ec2-user:~/hadoop-2.6.0/share$ cd ..
ubuntu@ec2-user:~/hadoop-2.6.0$ ls
bin  include  libexec  NOTICE.txt  sbin
etc  lib      LICENSE.txt  README.txt  share
ubuntu@ec2-user:~/hadoop-2.6.0$ cd sbin
ubuntu@ec2-user:~/hadoop-2.6.0/sbin$ ls
distribute-exclude.sh  start-all.cmd          stop-balancer.sh
hadoop-daemon.sh       start-all.sh           stop-dfs.cmd
hadoop-daemons.sh     start-balancer.sh       stop-dfs.sh
hdfs-config.cmd        start-dfs.cmd           stop-secure-dns.sh
hdfs-config.sh         start-dfs.sh            stop-yarn.cmd
httpfs.sh              start-secure-dns.sh     stop-yarn.sh
kms.sh                 start-yarn.cmd          yarn-daemon.sh
mr-jobhistory-daemon.sh start-yarn.sh           yarn-daemons.sh
refresh-namenodes.sh   stop-all.cmd
slaves.sh              stop-all.sh
ubuntu@ec2-user:~/hadoop-2.6.0/sbin$
```



```
ubuntu@ec2-user: ~/hadoop-2.6.0/etc/hadoop
RX bytes:11840 (11.8 KB) TX bytes:11840 (11.8 KB)

ubuntu@ec2-user:~$ cd etc
-bash: cd: etc: No such file or directory
ubuntu@ec2-user:~$ cd hadoop-2.6.0
ubuntu@ec2-user:~/hadoop-2.6.0$ ls
bin  include  libexec  NOTICE.txt  sbin
etc  lib      LICENSE.txt  README.txt  share
ubuntu@ec2-user:~/hadoop-2.6.0$ cd etc
ubuntu@ec2-user:~/hadoop-2.6.0/etc$ ls
hadoop
ubuntu@ec2-user:~/hadoop-2.6.0/etc$ cd hadoop
ubuntu@ec2-user:~/hadoop-2.6.0/etc/hadoop$ ls
capacity-scheduler.xml  httpfs-env.sh  mapred-env.sh
configuration.xml      httpfs-log4j.properties  mapred-queues.xml.template
container-executor.cfg  httpfs-signature.secret  mapred-site.xml.template
core-site.xml           httpfs-site.xml  slaves
hadoop-env.cmd          kms-acls.xml     ssl-client.xml.example
hadoop-env.sh           kms-env.sh       ssl-server.xml.example
hadoop-metrics2.properties  kms-log4j.properties  yarn-env.cmd
hadoop-metrics.properties  kms-site.xml       yarn-env.sh
hadoop-policy.xml        log4j.properties   yarn-site.xml
hdfs-site.xml            mapred-env.cmd
ubuntu@ec2-user:~/hadoop-2.6.0/etc/hadoop$
```

Cd etc/Hadoop

Set java home in Hadoop-env.sh

Command : sudo vi Hadoop-env.sh

And add following as in pic

```
ubuntu@ec2-user: ~/hadoop-2.6.0/etc/hadoop
# "License"); you may not use this file except in compliance
# with the License.  You may obtain a copy of the License at
#
#     http://www.apache.org/licenses/LICENSE-2.0
#
# Unless required by applicable law or agreed to in writing, software
# distributed under the License is distributed on an "AS IS" BASIS,
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
# See the License for the specific language governing permissions and
# limitations under the License.

# Set Hadoop-specific environment variables here.

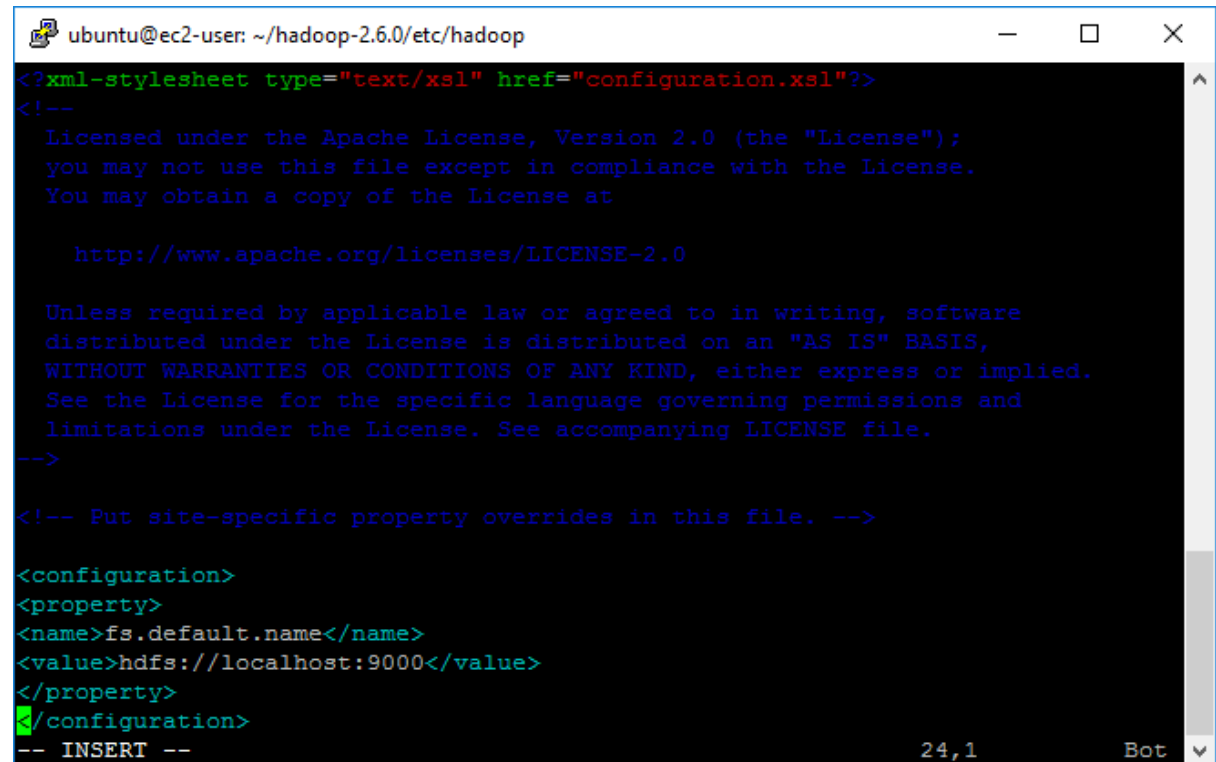
# The only required environment variable is JAVA_HOME.  All others are
# optional.  When running a distributed configuration it is best to
# set JAVA_HOME in this file, so that it is correctly defined on
# remote nodes.

# The java implementation to use.
export JAVA_HOME=${JAVA_HOME}
# Set JAVA_HOME

export JAVA_HOME=/usr/lib/jvm/java-9-openjdk-amd64/jre
:wq
```

Edit core-site.xml file. This file contains the configuration settings for Hadoop Core such as I/O settings that are common to HDFS and MapReduce.

Command: vi core-site.xml



```
ubuntu@ec2-user: ~/hadoop-2.6.0/etc/hadoop
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
<!--
Licensed under the Apache License, Version 2.0 (the "License");
you may not use this file except in compliance with the License.
You may obtain a copy of the License at

    http://www.apache.org/licenses/LICENSE-2.0

Unless required by applicable law or agreed to in writing, software
distributed under the License is distributed on an "AS IS" BASIS,
WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
See the License for the specific language governing permissions and
limitations under the License. See accompanying LICENSE file.
-->

<!--Put site-specific property overrides in this file. -->

<configuration>
<property>
<name>fs.default.name</name>
<value>hdfs://localhost:9000</value>
</property>
</configuration>
-- INSERT --
```

Command: vi hdfs-site.xml

```
ubuntu@ec2-user: ~/hadoop-2.6.0/etc/hadoop
Unless required by applicable law or agreed to in writing, software
distributed under the License is distributed on an "AS IS" BASIS,
WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
See the License for the specific language governing permissions and
limitations under the License. See accompanying LICENSE file.
-->

<!-- Put site-specific property overrides in this file. -->

<configuration>
<property>
<name>dfs.replication</name>
<value>1</value>
</property>
<property>
<name>dfs.namenode.name.dir</name>
<value>/home/ec2-user/hadoop-2.6.0/hdfs/namenode</value>
</property>
<property>
<name>dfs.datanode.data.dir</name>
<value>/home/ec2-user/hadoop-2.6.0/hdfs/datanode</value>
</property>
</configuration>
-- INSERT --
```

Edit mapred-site.xml. This file contains the configuration settings for MapReduce daemons.

Command: cp mapred-site.xml.template mapred-site.xml

Command: vi mapred-site.xml

```
ubuntu@ec2-user: ~/hadoop-2.6.0/etc/hadoop
<?xml version="1.0"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
<!--
  Licensed under the Apache License, Version 2.0 (the "License");
  you may not use this file except in compliance with the License.
  You may obtain a copy of the License at

    http://www.apache.org/licenses/LICENSE-2.0

  Unless required by applicable law or agreed to in writing, software
  distributed under the License is distributed on an "AS IS" BASIS,
  WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
  See the License for the specific language governing permissions and
  limitations under the License. See accompanying LICENSE file.
-->

<!-- Put site-specific property overrides in this file. -->

<configuration>
<property>
<name>mapreduce.framework.name</name>
<value>yarn</value>
</property>
</configuration>
~
~
-- INSERT --
```

Edit yarn-site.xml. This file contains the configuration settings for YARN.

Command: vi yarn-site.xml

```
ubuntu@ec2-user: ~/hadoop-2.6.0/etc/hadoop
<?xml version="1.0"?>
<!--
Licensed under the Apache License, Version 2.0 (the "License");
you may not use this file except in compliance with the License.
You may obtain a copy of the License at

    http://www.apache.org/licenses/LICENSE-2.0

Unless required by applicable law or agreed to in writing, software
distributed under the License is distributed on an "AS IS" BASIS,
WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
See the License for the specific language governing permissions and
limitations under the License. See accompanying LICENSE file.
-->
<!-- Site specific YARN configuration properties -->
<configuration>
<property>
  <name>yarn.nodemanager.aux-services</name> <value>mapreduce_shuffle</value>
</property>
<property>
  <name>yarn.nodemanager.aux-services.mapreduce.shuffle.class</name>
  <value>org.apache.hadoop.mapred.ShuffleHandler</value>
</property>
</configuration>

-- INSERT --
```

Command: cd

Command: hadoop namenode -format

Command: cd hadoop-2.6.0/sbin/

Command: ./start-dfs.sh

Start the YARN services i.e. ResourceManager and NodeManager

Command: ./start-yarn.sh

Now run jps command to check if the daemons are running.

Command: jps

Now open browser and enter public dns name and hit enter you will get a page running.

